

### PacketStream

(a platform for analyzing network traffic in near real-time)

#### Leonid Isaev

Insight Data Eng. Fellowship NY 2019



#### Goals:

- Build a distributed platform for troubleshooting network failures;
- Make it understand Wireshark-like rules.

#### **Constraints:**

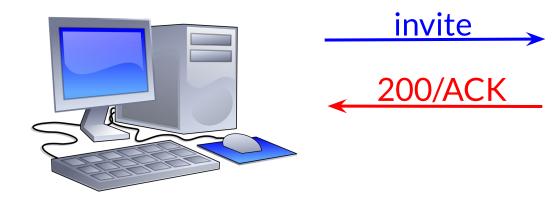
- Must be field-ready
  - A data producer is a CentOS/Ubuntu server;
  - No "pip install < 200 packages>" on clients.

# PacketStream architecture ubuntu analysis

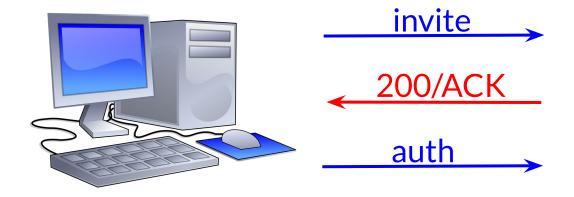


invite

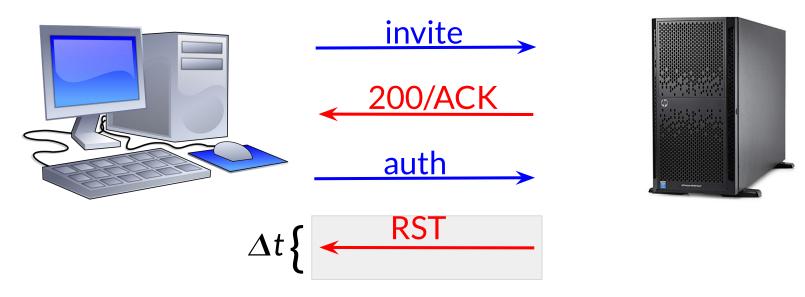








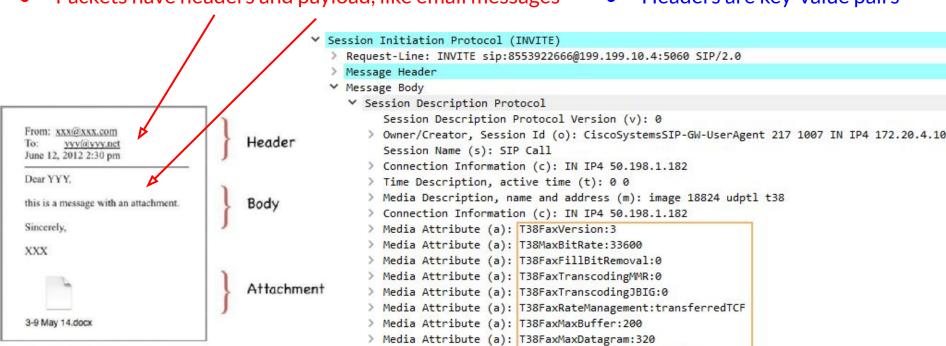




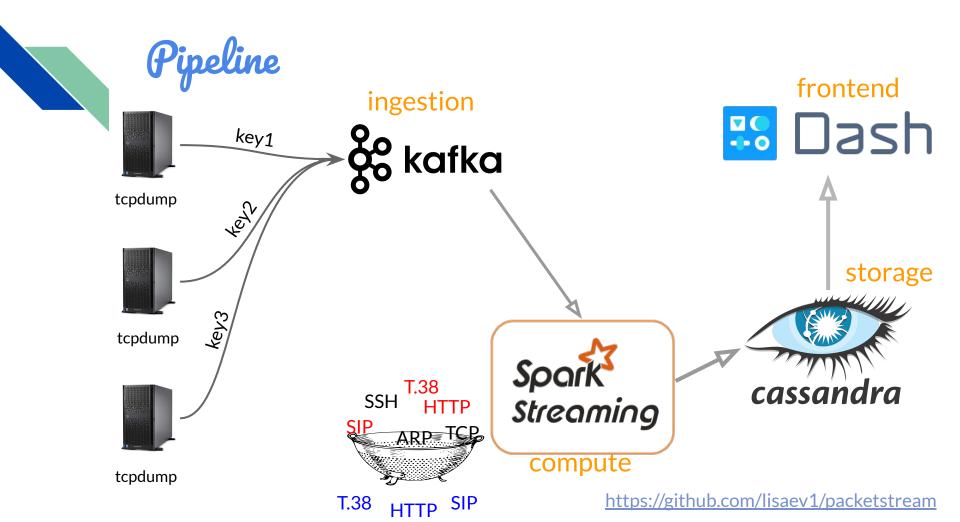
can't RST connection during wait window

## Sample packet structure (T.38 fax protocol)

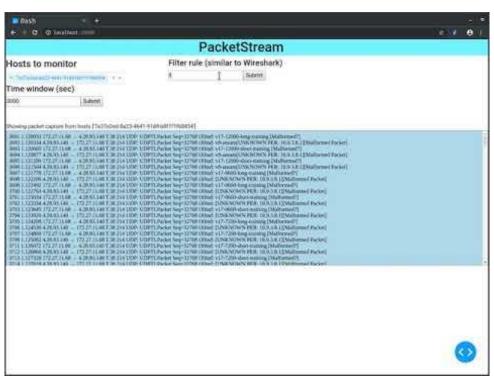
- Packets have headers and payload, like email messages
  - Headers are key-value pairs



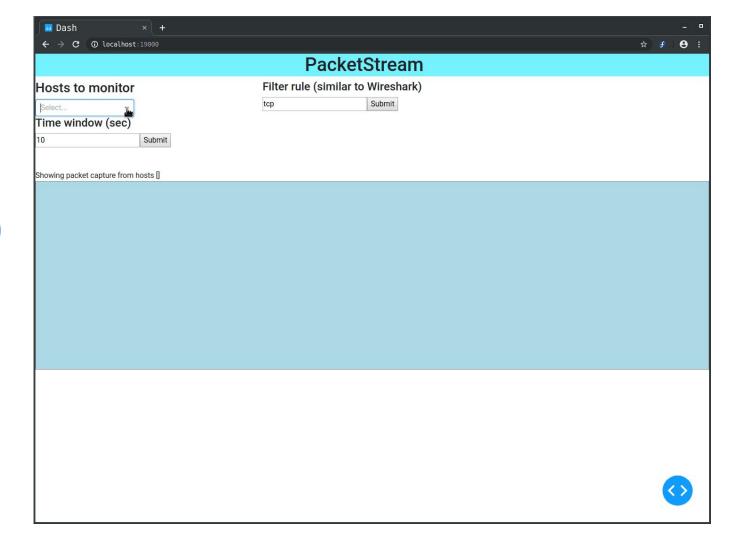
Media Attribute (a): T38FaxUdpEC:t38UDPRedundancy



### **Demo** (PacketStream)



# **Demo** (fallback)



# Challenges (messages with \$'\n' in Kafka)

**Problem:** Kafka treats '\n' as message separator

```
\xff\x89\x09\nsip:test@192\nX
intended
1 message

Kafka
topic

\xff\x89\x09

\sip:test@192

got 3!
```

### **Solution:** Use hexadecimal representation of the Python bytes objects

# Challenges (messages with \$'\n' in Kafka)

**Problem:** Kafka treats '\n' as message separator

- It's not possible with kafka-console-producer as it uses a Java Scanner object that's newline delimited.
- 2 You would need to do it via your own producer code
- share improve this answer

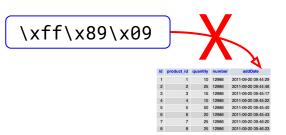
answered Sep 3 '18 at 21:27

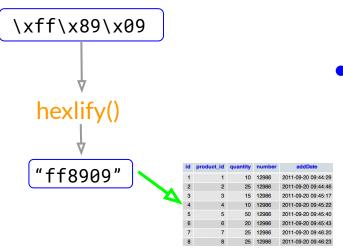
cricket\_007

92.4k • 11 • 58 • 139

## **Challenges** (writing network packets to Cassandra from PySpark)

- Spark DataFrame doesn't support Python bytes object
  - Network packets look like b"\xff\x89\x09...";
  - Need to convert RDD -> DF for writing to Cassandra.





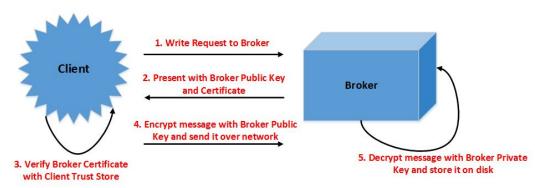
- Solution: store hex-encoded bytes
  - Base64 encoding can be wasteful;
  - Use binascii.hexlify(): b"\xff\x89\x09" -> "ff8909"

# **Challenges** (encrypted traffic from producers)

- Initial approach: wrap Kafka traffic in SSH tunnel
  - Works, but Kafka broker readvertisement doesn't happen
  - Slow -- SSH limits packet length (compared to http/https)

- Solution: use native Kafka TLS protocol
  - Certs are a PAIN
  - CA vs self-signed





#### Short bio

#### Background

- Indiana Univ Bloomington (computational physics)
- Postdoc at LSU Baton Rouge, Los Alamos, JILA/NIST/CU Boulder
- Linux support/software engineer @ iFax Solutions, Norristown PA









# FAX™ Solutions, Inc. TELEPHONY DEPOT 215-825-8700



#### Fun facts/hobbies:

- Real-life Sheldon
- Science-fiction, hacking on Arch Linux

### Appendix

### Kafka topic structure



9f1faa88-fa44-4ce9-b2ab-eb267630cebf

 $\longrightarrow \hspace{0.2in} \triangleright$ 

packets: bytestream



```
message = {
     "key": 9f1faa88-fa44-4ce9-b2ab-eb267630cebf,
     "value": b"..."
}
```

- Assign a UUID to each machine
  - Can be license number;
  - Python uuid.uuid4();
- UUID -> key in Kafka topic;
- Choose # of partitions = # of nodes in the cluster
  - Packets from each host end up in the same partition!

#### Cassandra table structure

hostid (uuid)	pktid (timeuuid)	payload (text)
7e37e2ed-8a23-4641-91d9-b8f1f1f68454	6cde8df0-f019-11e9-b24e-52a33b5da58a	0050569e7115b
89eab76b-2969-4031-9f27-7e53614edc55	6c6cd4ee-f019-11e9-b24e-52a33b5da58a	0050716e294b6
a704ae05-032f-426c-a211-e6249e885287	6c6d31a0-f019-11e9-b24e-52a33b5da58a	0050af98223c6

### Examples of broken networks

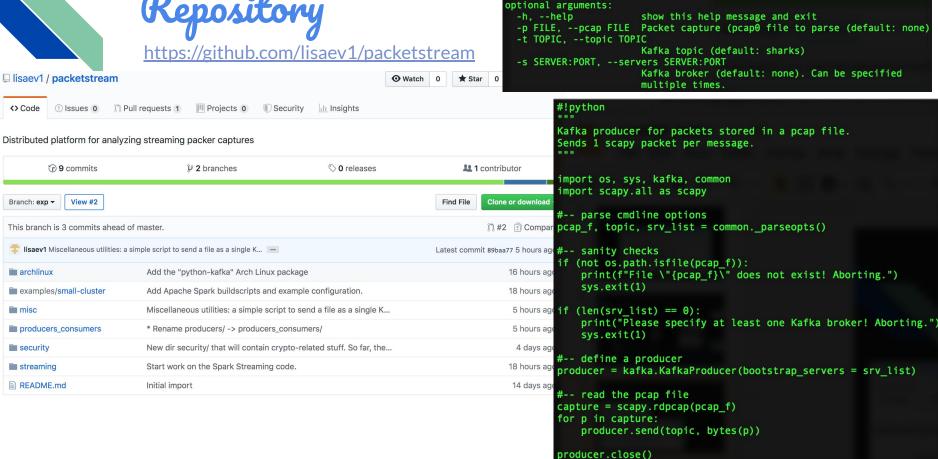
No	Time	Source	Destination	Protocol	Info
	1 0.000000	ipphone	Broadcast	ARP	Who has 192.168.1.10? Tell 192.168.1.151
	2 0.052480	pbxhost	ipphone	ARP	192.168.1.10 is at 00:10:5a:e1:90:6e
	3 0.000146	ipphone	pbxhost	SIP/SDP	Request: INVITE sip:sip%3a613@fwd.pulver.com, with session description
	4 0.001480	pbxhost	ipphone	SIP	Status: 407 Proxy Authentication Required
	5 0.052920	ipphone	pbxhost	SIP	Request: ACK sip:sip%3a613@fwd.pulver.com
	6 0.062692	ipphone	pbxhost	SIP/SDP	Request: INVITE sip:sip%3a613@fwd.pulver.com, with session description
	7 0.063859	pbxhost	ipphone	SIP	Status: 407 Proxy Authentication Required
	8 0.123275	ipphone	pbxhost	SIP	Request: ACK sip:sip%3a613@fwd.pulver.com
10	9 0.132364	ipphone	pbxhost	SIP/SDP	Request: INVITE sip:sip%3a613@fwd.pulver.com, with session description
1	0 0.135772	pbxhost	ipphone	SIP	Status: 407 Proxy Authentication Required
1	1 0.193816	ipphone	pbxhost	SIP	Request: ACK sip:sip%3a613@fwd.pulver.com
1	2 0.202913	ipphone	pbxhost	SIP/SDP	Request: INVITE sip:sip%3a613@fwd.pulver.com, with session description
1	3 0.208640	pbxhost	ipphone	SIP	Status: 407 Proxy Authentication Required
1	4 0.258099	ipphone	pbxhost	SIP	Request: ACK sip:sip%3a613@fwd.pulver.com

- Broken firmware -> timing issues, random connection resets;
- Wrong credentials, encoding, etc.

## Packet structure (broken packets)

```
Frame 611: 214 bytes on wire (1712 bits), 214 bytes captured (1712 bits)
Ethernet II, Src: Fortinet_09:00:04 (00:09:0f:09:00:04), Dst: Vmware_9e:71:15 (00:50:56:9e:71:15)
Internet Protocol Version 4, Src: 4.28.93.140, Dst: 172.27.11.68
User Datagram Protocol, Src Port: 20458, Dst Port: 56008
ITU-T Recommendation T.38
    [Stream setup by SDP (frame 598)]
        [Stream frame: 598]
        [Stream Method: SDP]
    UDPTLPacket
        seq-number: 32768
        something unknown here [10.9 Unconstrained unexpected fragment count]
            [Expert Info (Warning/Undecoded): something unknown here [10.9 Unconstrained unexpected fragment count]]
                 [something unknown here [10.9 Unconstrained unexpected fragment count]]
                [Severity level: Warning]
                [Group: Undecoded]
Malformed Packet: T.38
```





usage: scapy\_prod.py [-h] [-p FILE] [-t TOPIC] [-s SERVER:PORT]

### Test run -- a look inside Spark Stream

CO2RJ1JPG8WM:packetstream insight\$ python3 producers\_consumers/python/scapy\_prod.py -p ~/Downloads/\_] ifax.pcap -t "test1" \$(for i in {1..3}; do echo -nE "-s 192.168.122.7\${i}:9092 "; done)

```
lisaev@spark-node2 ~ spark-submit --master spark://spark-node1:7077 --jars ks.jar ./echo.py 192.1
68.122.71:9092 test1
<bound method Packet.summary of <Ether    dst=00:09:0f:09:00:04 src=00:50:56:9e:71:15 type=IPv4 |<IP</pre>
version=4 ihl=5 tos=0x0 len=931 id=28856 flags=DF frag=0 ttl=64 proto=udp chksum=0xad8a src=172.27.1
1.68 dst=4.28.93.140 | <UDP | sport=sip dport=sip len=911 chksum=0x1ca8 | <Raw | load='INVITE | sip:0ae9fa
eaad38d22c779775e03bb9d755@4.28.93.140:5060 SIP/2.0\r\nFrom: <sip:4012830857@sbc01-stl.sip.univergeb
lue.com>:tag=dcfef6b0-0-13c4-65014-a5c-75d5ed11-a5c\r\nTo: "PHILADELPHIA PA"<sip:2153830199@sbc01-st
l.sip.univergeblue.com>;tag=4.28.93.140+1+3218a8f7+270ff97\r\nCall-ID: 0g0AAC8WAAACBAAALxYAANYw/3hJv
01vNg45pv0sNQbJaGaS1oYg5/3tEuM5AjEX@4.28.93.140\r\nCSeg: 1 INVITE\r\nVia: SIP/2.0/UDP 172.27.11.68:5
060;branch=z9hG4bK-a60-28871c-296f3a47-dd2ff5e8\r\nSupported: 100rel\r\nMax-Forwards: 70\r\nUser-Age
nt: Brktsip/6.9.6B10 (Dialogic)\r\nContact: <sip:172.27.11.68>\r\nContent-Type: application/sdp\r\nC
ontent-Length: 286\r\n\r\nv=0\r\no=- 2208997233 0083442001 IN IP4 172.27.11.68\r\ns=no session name\
r\nt=0 0\r\nm=image 56008 udptl t38\r\nc=IN IP4 172.27.11.68\r\na=T38FaxVersion:0\r\na=T38MaxBitRate
:14400\r\na=T38FaxRateManagement:transferredTCF\r\na=T38FaxMaxBuffer:200\r\na=T38FaxMaxDatagram:72\r
\na=T38FaxUdpEC:t38UDPRedundancv\r\n' |>>>>
<bound method Packet.summary of <Ether dst=00:50:56:9e:71:15 src=00:09:0f:09:00:04 type=IPv4 |<IP</pre>
version=4 ihl=5 tos=0x0 len=491 id=49541 flags= frag=0 ttl=64 proto=udp chksum=0x9e75 src=4.28.93.14
0 dst=172.27.11.68 | <UDP | sport=sip dport=sip len=471 chksum=0x9451 | <Raw | load='SIP/2.0 100 Trying\
r\nCall-ID: 0g0AAC8WAAACBAAALxYAANYw/3hJv01vNq45pv0sN0bJaGaS1oYg5/3tEuM5AjEX@4.28.93.140\r\nCSeq: 1
INVITE\r\nFrom: <sip:4012830857@sbc01-stl.sip.univergeblue.com>;tag=dcfef6b0-0-13c4-65014-a5c-75d5ed
11-a5c\r\nTo: "PHILADELPHIA PA"<sip:2153830199@sbc01-stl.sip.univergeblue.com>;tag=4.28.93.140+1+321
8a8f7+270ff97\r\nVia: SIP/2.0/UDP 172.27.11.68:5060:branch=z9hG4bK-a60-28871c-296f3a47-dd2ff5e8\r\nS
erver: SIP/2.0\r\nContent-Length: 0\r\nUser-Agent: SoTel\r\n\r\n' |>>>>
```

### Roadmap: putting all this into production

Nobody should start to undertake a large project. You start with a small *trivial* project, and you should never expect it to get large. If you do, you'll just overdesign and generally think it is more important than it likely is at that stage. So start small, and think about the details. Don't think about some big picture and fancy design. If it doesn't solve some fairly immediate need, it's almost certainly over-designed.

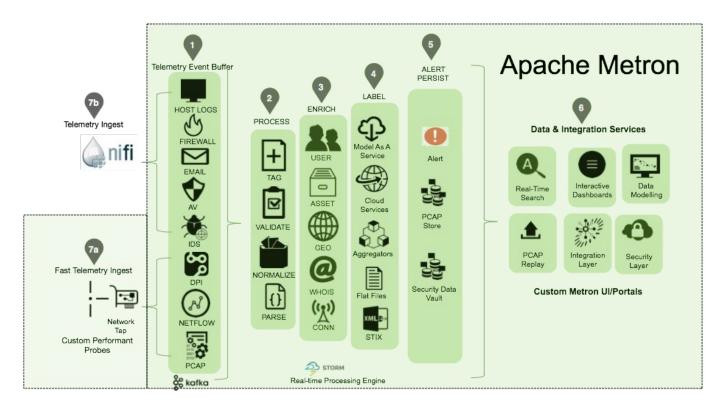
-- Linus Torvalds, 2004

- CLI interface (Dash is great, but non-functional):
  - Need to look at raw packets;
  - Wireshark rules/filters;
- Richer processing rules in Spark, incl wireshark rule files;
- Dynamic provisioning of Kafka brokers and Spark nodes;
- Dedicated distributed storage (perhaps outside Cassandra).

complexity

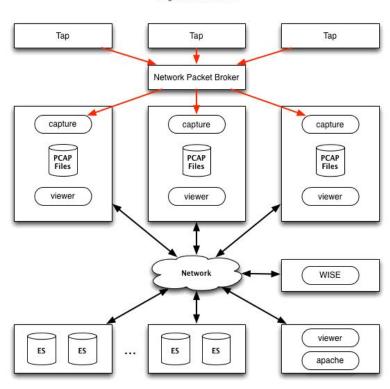


### Related work: Apache Metron



#### Related Work: AOL Moloch

Sample High Traffic Deployment Single Moloch Cluster



### Why not use Kafka streaming?

- Streaming apps have to be started manually on each broker with
   # of instances = # of partitions in the consumed topic;
- No automatic re-spawning in case of node failure;
- Complex code (lots of async stuff);
- Processing and message handling are co-located.

### We did test Faust streaming for Kafka



#### **Faust**

A library for building streaming applications in Python.



build failing

#### Navigation

Copyright

Introducing Faust

Playbooks

User Guide

Frequently Asked Questions (FAQ)

API Reference

Changes

Contributing

Developer Guide

History

#### Faust - Python Stream Processing

```
# Python Streams ٩(@_@)?
# Forever scalable event processing & in-memory durable K/V store;
# w/ asyncio & static typing.
import faust
```

**Faust** is a stream processing library, porting the ideas from Kafka Streams to Python.

It is used at <u>Robinhood</u> to build high performance distributed systems and real-time data pipelines that process billions of events every day.

Faust provides both stream processing and event processing, sharing similarity with tools such as Kafka Streams, Apache Spark/Storm/Samza/Flink,

It does not use a DSL, it's just Python! This means you can use all your favorite Python libraries when stream processing: NumPy, PyTorch, Pandas, NLTK, Django, Flask, SQLAlchemy, ++

Faust requires Python 3.6 or later for the new <u>async/await</u> syntax, and variable type annotations.

Here's an example processing a stream of incoming orders:

```
app = faust App('myapp', broker='kafka://localhost')

# Models describe how messages are serialized:
# {"account_id": "3fae-...", amount": 3}

class Order(faust Record):
    account_id: str
    amount: int

@app.agent(value_type=Order)
async def order(orders):
    async for order in orders:
```